HIVE

Es una infraestructura para el almacenaje y consultade datos basada en Hadoop

Hadoop proporciona escalabilidad masiva y con capacidades de tolerancia a fallos para el procesamiento y almacenamiento de datos

Ejercicios Hive

1. Entrar en Hive

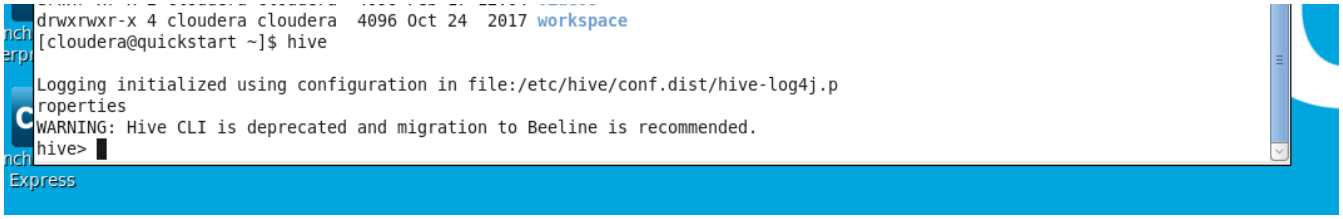

2. Modificar la propiedad correspondiente para mostrar por pantalla las cabeceras de las tablas

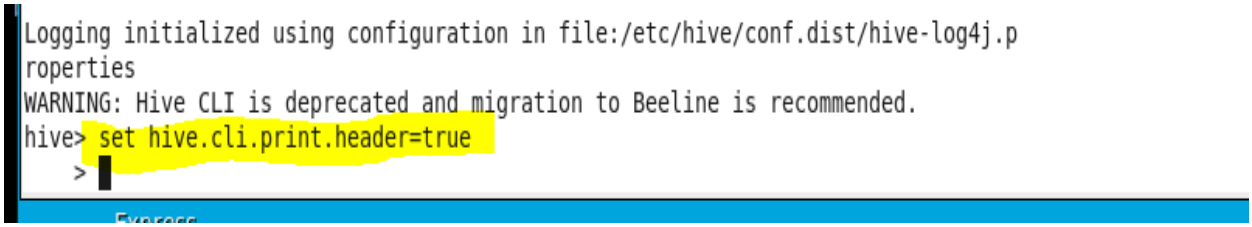

3. Crear una base de datos llamada “cursohivedb”

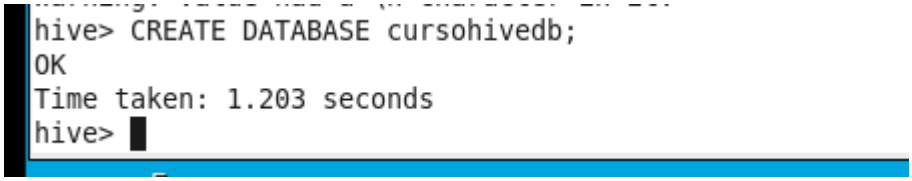

4. Situarnos en la base de datos recién creada para trabajar con ella

5. Comprobar que la base de datos está vacía

6. Crear una tabla llamada “iris” en nuestra base de datos que contenga 5 columnas

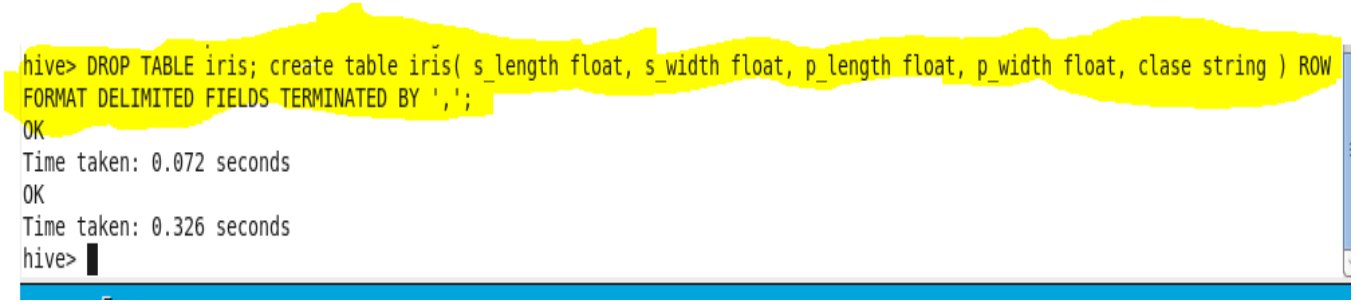

7. Comprobar que la tabla se ha creado y el tipado de sus columnas

```
hive> SHOW TABLES;
OK
iris
Time taken: 0.129 seconds, Fetched: 1 row(s)
hive>
```

Express

Tipado de sus columnas

```
hive> DESC iris;
OK
s_length                float
s_width                 float
p_length                float
p_width                 float
clase                   string
Time taken: 0.128 seconds, Fetched: 5 row(s)
hive>
```

8. Importar el fichero “iris_completo.txt” al local file system del cluster en la carpeta /home/cloudera/ejercicios/ejercicios_HIVE

a. Copiar el fichero a HDFS en la ruta /user/cloudera/hive. Realizar las acciones Necesarias

```
hive> DROP TABLE iris; create table iris( s_length float, s_width float, p_length float, p_width float, clase string
) ROW FORMAT DELIMITED FIELDS TERMINATED BY ','STORED AS TEXTFILE LOCATION '/home/cloudera/ejercicios/ejercicios_HIVE
/iris_completo.tx';
OK
Time taken: 0.072 seconds
OK
Time taken: 0.15 seconds
hive>
```

9. Comprueba que el fichero está en la ruta en HDFS indicada

```
[cloudera@quickstart ~]$ mkdir hive
[cloudera@quickstart ~]$ ls
cloudera-manager  eclipse                     kerberos  Public
cm_api.py         ejercicios                  lib       Templates
Desktop           enterprise-deployment.json  Music     Videos
Documents         express-deployment.json     parcels   workspace
Downloads         hive                        Pictures
[cloudera@quickstart ~]$
```

```
[cloudera@quickstart /]$ hadoop fs -put /home/cloudera/ejercicios/ejercicios_HIV
E/iris_completo.txt /user/cloudera/hive
[cloudera@quickstart /]$ ls
```

10. Importa el fichero en la tabla iris que acabamos de crear desde HDFS

```
[cloudera@quickstart /]$ hadoop fs -put /home/cloudera/ejercicios/ejercicios_HIVE/iris_com
pleto.txt /user/cloudera/hive
put: `/user/cloudera/hive/iris_completo.txt': File exists
```

11. Comprobar que la table tiene datos

```
hive> SELECT * FROM iris;
OK
1.0     3.2     4.3     5.7     Iris-virginica
Time taken: 0.049 seconds, Fetched: 1 row(s)
hive>
```

12. Mostrar las 5 primeras filas de la tabla iris

```
Hive> SELECT * FROM iris Limit 5;
```

13. Mostrar solo aquellas filas cuyo s_length sea mayor que 5. Observad que se ejecuta un MapReduce y que el tiempo de ejecución es un poco mayor

```
Hive> Select * from iris as i where i.s_length>5
```

14. Seleccionar la media de s_width agrupados por clase. Observad que ahora el tiempo de ejecución aumenta considerablemente.

15. Pregunta: vemos que aparece un valor NULL como resultado en la query anterior. ¿Por qué? ¿cómo los eliminarías? Porque había algún dato erróneo, no numérico o nulo en el campo de

alguna clase. Para eliminarlos podríamos añadir la condición where para que fuera distinto de null.

16. Insertar en la tabla la siguiente fila (1.0,3.2,4.3,5.7,"Iris-virginica")

```
hive> insert into table iris values (1.0,3.2,4.3,5.7,"Iris-virginica");
Query ID = cloudera_20220406164545_1e2368d7-4a64-4ff7-b5b8-329d5be55f2c
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1649143967398_0001, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1649143967398
_0001/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1649143967398_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2022-04-06 16:45:42,210 Stage-1 map = 0%,  reduce = 0%
2022-04-06 16:45:48,892 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.19 sec
MapReduce Total cumulative CPU time: 1 seconds 190 msec
Ended Job = job_1649143967398_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://quickstart.cloudera:8020/home/cloudera/ejercicios/ejercicios_HIVE/iris_completo.tx/.hive-stagi
ng_hive_2022-04-06_16-45-32_321_9093019405324985248-1/-ext-10000
Loading data to table default.iris
Table default.iris stats: [numFiles=2, numRows=1, totalSize=31, rawDataSize=30]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 1.19 sec   HDFS Read: 4788 HDFS Write: 99 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 190 msec
OK
Time taken: 17.883 seconds
hive> SELECT * FROM iris;
OK
1.0     3.2     4.3     5.7     Iris-virginica
Time taken: 0.049 seconds, Fetched: 1 row(s)
hive>
```

17. Contar el número de ocurrencias de cada clase

18. Seleccionar las clases que tengan más de 45 ocurrencias a. Select clase from iris group by clase having count(*)>45;

```
hive> select clase, count(*)
    >
    > from iris
    >
    > group by clase;
Query ID = cloudera_20220406170505_b65fa69c-97b0-403e-b204
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input
In order to change the average load for a reducer (in byte
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1649143967398_0004, Tracking URL = http
_0004/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1
Hadoop job information for Stage-1: number of mappers: 2;
2022-04-06 17:05:34,625 Stage-1 map = 0%,  reduce = 0%
2022-04-06 17:05:44,424 Stage-1 map = 100%,  reduce = 0%,
2022-04-06 17:05:52,793 Stage-1 map = 100%,  reduce = 100%
MapReduce Total cumulative CPU time: 2 seconds 820 msec
Ended Job = job_1649143967398_0004
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2  Reduce: 1   Cumulative CPU: 2.82 se
Total MapReduce CPU Time Spent: 2 seconds 820 msec
OK
Iris-virginica  2
Time taken: 26.15 seconds, Fetched: 1 row(s)
```

19. Utilizando la función LEAD, ejecutar una query que devuelva la clase, p_length y el LEAD de p_length con Offset=1 y Default_Value =0, particionado por clase y ordenado por p_length.

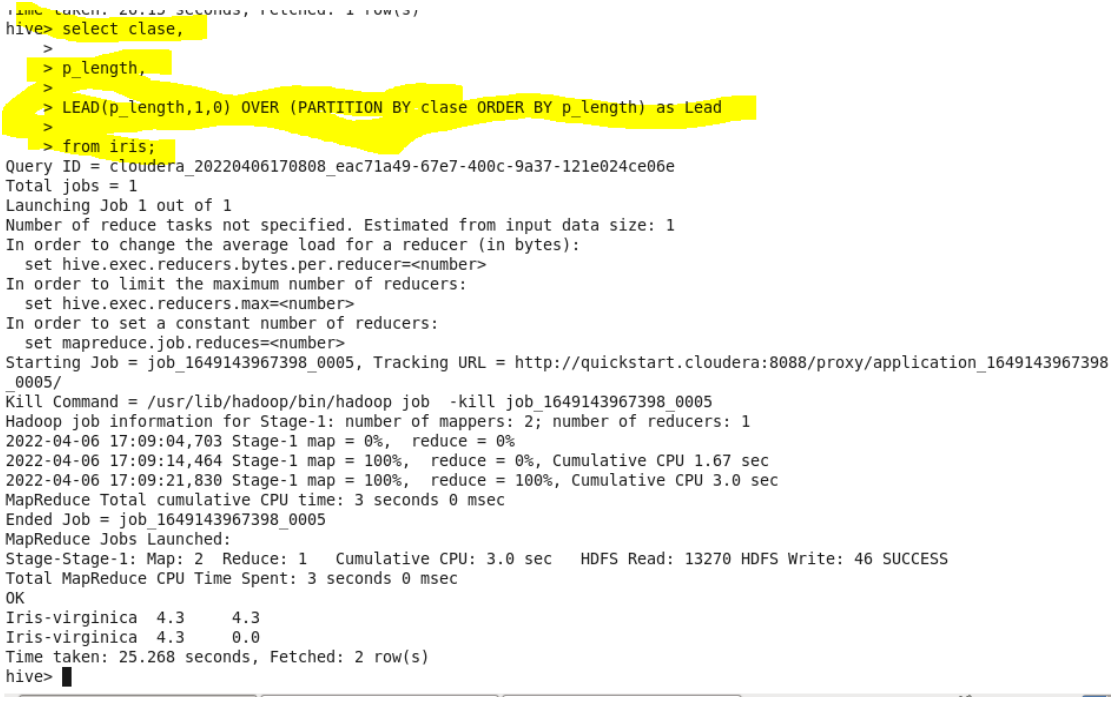

```
hive> select clase,
    >
    > p_length,
    >
    > LEAD(p_length,1,0) OVER (PARTITION BY clase ORDER BY p_length) as Lead
    >
    > from iris;
Query ID = cloudera_20220406170808_eac71a49-67e7-400c-9a37-121e024ce06e
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1649143967398_0005, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1649143967398
_0005/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1649143967398_0005
Hadoop job information for Stage-1: number of mappers: 2; number of reducers: 1
2022-04-06 17:09:04,703 Stage-1 map = 0%,  reduce = 0%
2022-04-06 17:09:14,464 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.67 sec
2022-04-06 17:09:21,830 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 3.0 sec
MapReduce Total cumulative CPU time: 3 seconds 0 msec
Ended Job = job_1649143967398_0005
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2  Reduce: 1   Cumulative CPU: 3.0 sec   HDFS Read: 13270 HDFS Write: 46 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 0 msec
OK
Iris-virginica  4.3     4.3
Iris-virginica  4.3     0.0
Time taken: 25.268 seconds, Fetched: 2 row(s)
hive>
```

20. Utilizando funciones de ventanas, seleccionar la clase, p_length, s_length, p_width, el número de valores distintos de p_length en todo el dataset, el valor máximo de s_length por clase y la media de p_width por clase, ordenado por clase y s_length de manera descendente.

```
hive> select clase,
    >
    > p_length,
    >
    > s_length,
    >
    > p_width,
    >
    > count(p_length) over (partition by p_length) as pl_ct,
    >
    > max(s_length) over (partition by clase) as sl_ct,
    >
    > avg(p_width) over (partition by clase) as sl_av
    >
    > from iris
    >
    > order by clase,s_length desc;
Query ID = cloudera_20220406171212_a5873796-b686-4e3d-8f23-13b6e82b564f
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1649143967398_0006, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1649143967398
_0006/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1649143967398_0006
Hadoop job information for Stage-1: number of mappers: 2; number of reducers: 1
2022-04-06 17:12:30,168 Stage-1 map = 0%,  reduce = 0%
2022-04-06 17:12:38,781 Stage-1 map = 50%,  reduce = 0%, Cumulative CPU 0.88 sec
2022-04-06 17:12:39,807 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.67 sec
2022-04-06 17:12:46,192 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 2.89 sec
MapReduce Total cumulative CPU time: 2 seconds 890 msec
Ended Job = job_1649143967398_0006
Launching Job 2 out of 3
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1649143967398_0007, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1649143967398
_0007/
```


```
Starting Job = job_1649143967398_0007, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1649143967398
_0007/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1649143967398_0007
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2022-04-06 17:12:53,931 Stage-2 map = 0%,  reduce = 0%
2022-04-06 17:12:59,227 Stage-2 map = 100%,  reduce = 0%, Cumulative CPU 0.66 sec
2022-04-06 17:13:05,544 Stage-2 map = 100%,  reduce = 100%, Cumulative CPU 1.85 sec
MapReduce Total cumulative CPU time: 1 seconds 850 msec
Ended Job = job_1649143967398_0007
Launching Job 3 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1649143967398_0008, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1649143967398
_0008/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1649143967398_0008
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 1
2022-04-06 17:13:12,521 Stage-3 map = 0%,  reduce = 0%
2022-04-06 17:13:18,836 Stage-3 map = 100%,  reduce = 0%, Cumulative CPU 0.72 sec
2022-04-06 17:13:27,299 Stage-3 map = 100%,  reduce = 100%, Cumulative CPU 1.79 sec
MapReduce Total cumulative CPU time: 1 seconds 790 msec
Ended Job = job_1649143967398_0008
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2  Reduce: 1   Cumulative CPU: 2.89 sec   HDFS Read: 12825 HDFS Write: 186 SUCCESS
Stage-Stage-2: Map: 1  Reduce: 1   Cumulative CPU: 1.85 sec   HDFS Read: 8562 HDFS Write: 210 SUCCESS
Stage-Stage-3: Map: 1  Reduce: 1   Cumulative CPU: 1.79 sec   HDFS Read: 6853 HDFS Write: 102 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 530 msec
OK
Iris-virginica  4.3     1.0     5.7     2       1.0     5.699999809265137
Iris-virginica  4.3     1.0     5.7     2       1.0     5.699999809265137
Time taken: 65.02 seconds, Fetched: 2 row(s)
hive>
```